

Artificial Intelligence Governance: A

Personal Operating Framework

Justin Simmons · A Personal Operating Framework

Purpose

This framework states the principle by which I evaluate artificial intelligence and the practices that follow from it. It is written to be applied before a tool is adopted, not consulted after its outcomes have already reached the people it affects. It is intended for a mission-driven institutional context, where enthusiasm for new AI initiatives is constant and the constraints on institutional data are real.

Foundational Principle: AI Is a Contained Tool, Not a Neutral One

Tools range across a spectrum of bias, from the tool-agnostic to the bias-laden. Some tools are agnostic in their utility. A pencil lends no bias to its user: in the hands of a revolutionary, a dictator, or a student, it remains the same instrument and serves each equally. A sword sits at the other end. Though it can be turned to self-defense, it has no agnostic use, because its purpose, from its design to the intent of its wielder, is to harm another. Even the epee,

the sword refined into sporting equipment, makes the point. Only through strictly enforced rules and protective gear can a tool built to wound be made safe enough for sport. The bias was never removed. It was contained.

Artificial intelligence resists being placed at a single point on that spectrum, because it is not a point but a gradation in itself. Some bias is intrinsic to any human-supervised model, since even a simple labeling decision encodes human judgment during training. There is no truly agnostic AI model where a person defined what the right answer looked like. An unsupervised model escapes that particular source of bias, but it does not escape bias itself. It trades the bias of human labels for the bias buried in the data it learns from, and it does so with no one in the loop to notice when the patterns it finds are skewed. In that sense, unsupervised learning can be less overseen without being more impartial.

This matters because of what AI is becoming. More than perhaps any tool in human history, it has the capacity to shape not only what we decide but how we think. A pencil records a thought; AI can author one. That is precisely why its intrinsic bias can never be an afterthought. It must be understood, named, and held in view whenever the tool is used, and most of all in the moments that matter.

Everything below follows from treating AI as contained rather than neutral. If the bias cannot be removed, then governance is the work of locating it, naming it, and accounting for it before it acts.

Operating Practice 1: Interrogate Provenance Before Adoption

At a small, mission-driven institution, AI has worked its way into nearly every decision, and ideas for new AI initiatives arrive constantly. People are often surprised when I meet their enthusiasm with questions rather than agreement, but this is the first and most important safeguard, and it follows directly from seeing AI as a contained rather than a neutral tool.

Before adoption, I determine who is behind the tool, where the data will be stored, and how that data will be used. Any "free" tools are met with overt skepticism, because the data shared with them will almost certainly become training material, and once it does, you do not get it back. In personal life, feeding your data to someone else's model is already a problem. With institutional data, including records protected by FERPA or HIPAA, it is a serious one, and knowing where that data lives after it is processed is paramount.

Because a tool's data and its bias are inseparable, tracing that origin is also how I judge what bias the system carries. A model is only ever as representative as the data behind it, and I cannot account for a bias whose source I cannot see. The goal is not to slow good ideas down. It is to ensure the bias intrinsic to the tool is named and accounted for before it touches a student, rather than discovered afterward in its outcomes.

Operating Practice 2: Prefer On-Premises Deployment

Use on-premises solutions whenever possible. Running a model on institution-owned hardware does more than satisfy the FERPA and HIPAA concerns from the first practice. It is also what makes genuine bias vetting possible. When the model lives on our own infrastructure, I control which model is used, I can evaluate it against our own students before it is deployed, and it cannot be silently swapped or updated underneath us, which would render any prior vetting meaningless.

This matters because of the pace of the field. New models arrive almost weekly, frontier models are often deprecated as soon as their successors launch, and users are funneled into an early-adoption cycle that pushes them onto unverified systems before proper evaluation can happen. On-premises deployment breaks that funnel. We adopt a model when we have vetted it, not when a vendor decides to sunset the last one.

This does not mean cloud models should never be used. Frontier models too often are not available to run locally, and when you need the best tool available, turning to the cloud

makes sense. It should be done cautiously, however, and with the knowledge that data shared this way is most likely not FERPA or HIPAA compliant.

Operating Practice 3: Govern Against Automation Bias

AI is, for all its usefulness, sycophantic to a fault. For an impassioned user, a model will often fabricate data or make logical leaps to "help," making the user's own argument seem not just feasible but natural and well-supported. This is what makes automation bias so dangerous. It is not only that people over-trust a machine's output, but that the machine actively bends toward telling them what they want to hear. The flattery of the tool and the credulity of the user reinforce each other, manufacturing a false confidence neither would reach alone.

Guarding against this cannot be left to good intentions. Users must understand the failure mode, verify the model's claims against primary sources, and treat a single AI session as a draft rather than a verdict, cross-checking across models and conversations. Most importantly, ideas should be tested with human colleagues before they are pushed up the chain.

■ *The final safeguard against a machine that agrees with you is a person who will not.*