

Artificial Intelligence: A Personal Operating Framework

Justin Simmons · A Personal Operating Framework

Introduction

This framework states the principle by which I evaluate artificial intelligence and the practices that follow from it. It is written to be applied before a tool is adopted, not consulted after its outcomes have already reached the people it affects. It is intended for a mission-driven institutional context, where enthusiasm for new AI initiatives is constant and the constraints on institutional data are real.

The framework rests on a single foundational conviction: AI is a contained tool, not a neutral one. Tools range across a spectrum of bias, from the tool-agnostic to the bias-laden. A pencil lends no bias to its user: in the hands of a revolutionary, a dictator, or a student, it remains the same instrument and serves each equally. A sword sits at the other end. Though it can be turned to self-defense, it has no agnostic use, because its purpose, from its design to the intent of its wielder, is to harm another. Even the epee, the sword refined into sporting equipment, makes the point. Only through strictly enforced rules and protective gear can a tool built to wound be made safe enough for sport. The bias was never removed. It was contained.

Artificial intelligence resists being placed at a single point on that spectrum, because it is not a point but a gradation in itself. Some bias is intrinsic to any human-supervised model, since even a simple labeling decision encodes human judgment during training. An unsupervised model escapes that particular source of bias, but it does not escape bias itself: it trades the bias of human labels for the bias buried in the data it learns from, with no one in the loop to notice when the patterns it finds are skewed. This matters because of what AI is becoming. More than perhaps any tool in human history, it has the capacity to shape not only what we decide but how we think. A pencil records a thought; AI can author one.

Everything below follows from treating AI as contained rather than neutral. If the bias cannot be removed, then leadership in this field is the work of locating it, naming it, and accounting for it before it acts.

Mission Statement

To lead the adoption of artificial intelligence with a spirit of collaboration instead of competition, allowing AI to be used to enhance productivity and output of human workers rather than replacing them. Every tool is approached with curiosity first, asking how it can be used safely, effectively, and ethically. Colleagues are equipped to use AI critically rather than passively,

ensuring tools are adopted for a purpose, not simply because they are available. Lastly, human judgment remains the final safeguard in every decision AI informs.

Core Values

Stewardship of what is entrusted to us. Institutional data represents real people — students, colleagues, patients — who did not consent to becoming training material. Records protected by FERPA or HIPAA are a trust, not a resource, and knowing where data lives after it is processed is a non-negotiable obligation of leadership. This value aligns directly with the ethic of my institution: the people behind the data are the mission, and the technology exists to serve them, never the reverse.

Honesty about what the tool is. AI is not neutral, and pretending otherwise is the first ethical failure from which the others follow. Naming a system's bias and its limitation is the bedrock of any decision guided by AI, from its implementation to its everyday use. As AI integration becomes more commercialized, this honesty ensures the user and their workplace remain steadfast in their mission, not the mission of the AI company.

Human judgment as the final authority. AI is, for all its usefulness, sycophantic to a fault. The flattery of the tool and the credulity of the user reinforce each other, manufacturing a false confidence neither would reach alone. The final safeguard against a machine that agrees with you is a person who will not. I hold this value in both directions: I seek out colleagues who will challenge my AI-assisted conclusions, and I strive to be that colleague for others.

Deliberation over urgency. New models arrive almost weekly, and vendors funnel users into an early-adoption cycle that pushes them onto unverified systems before proper evaluation can happen. Enthusiasm and effective marketing are not a vetting process.

Key Objectives

1. **Complete the Master of Science in Artificial Intelligence and Data Analytics** at Indiana Wesleyan University by July 2027, maintaining strong academic standing through the remaining course sequence.
2. **Establish a formal AI vetting process** at my institution within the next twelve months: a documented review (who is behind the tool, where data is stored, how data will be used) applied to every proposed AI adoption before deployment. Users who choose tools outside of institutional choices are not opponents, but rather collaborators who inform future decisions.
3. **Deploy and evaluate at least one on-premises large language model** against institutional use cases within the next year, demonstrating that locally hosted models can satisfy privacy requirements that cloud tools cannot, and documenting the evaluation methodology so it is repeatable.

4. **Foster a culture of iterative learning** by mentoring colleagues in critical AI use both in the IT Services department as well as other staff and faculty areas. Success is measured by their independent application of these practices in their own work.
5. **Transition into a formally AI-focused role** within my institution as the practical validation of this framework, with the documented deployments and vetting processes above serving as the evidence base.

Action Plans

Short-term (0–12 months). Complete the current graduate course sequence, applying each course's material to a live institutional problem rather than treating coursework and work as separate tracks. Formalize the review checklist currently applied informally and propose it with institutional leadership, supported by documented outcomes rather than abstract argument. Continue hands-on evaluation of locally hosted models in a personal lab environment, translating into institution-appropriate proposals. Begin mentoring relationships across campus departments, giving mentees frameworks they can modify rather than procedures they must follow.

Long-term (1–3 years). Use the knowledge from the graduate program, along with hands-on experiences, to legitimize a formal AI leadership responsibility. Expand the vetting process from a personal practice into an institutional policy with cross-functional ownership, engaging colleagues in academic affairs, information security, and compliance so the framework survives beyond any single person. Pursue mentorship from practitioners who have led AI governance at larger institutions, and contribute back by presenting the institution's approach in regional higher-education technology communities.

Cross-functional collaboration. Because AI adoption decisions at a small institution touch every department, I will continue building relationships outside my own team so that governance conversations happen with stakeholders rather than at them. Transparency about the real cost of requests and initiatives is part of this practice: colleagues make better decisions when they know what a proposal actually requires.

Evaluation and Adaptation

Self-assessment schedule. At the end of each graduate course (roughly every six weeks), I will review this framework against three questions: Did I apply the vetting process to every AI decision I touched? Where did automation bias appear in my own work, and did I catch it or did someone else? Are my objectives still the right ones, or has the field shifted underneath them?

Feedback mechanisms. The framework's own third value supplies its evaluation method: ideas are tested with human colleagues before they are pushed up the chain. I will continue to seek direct, critical feedback from my manager, from the colleagues I mentor, and from peers outside

my department who have no stake in agreeing with me. AI-assisted self-review is permitted but never sufficient; a single AI session is a draft, not a verdict.

Adaptation triggers. Some revisions should not wait for a scheduled review. A significant regulatory change, a data-handling incident anywhere in higher education relevant to our stack, or the deprecation of a vetted model are each immediate triggers to reopen the framework. The pace of this field is itself one of the framework's core concerns, and the framework must be held to the same standard it applies to the tools it governs: vetted, revisited, and never silently out of date.

Continuous improvement. This document is versioned deliberately. Each revision will record what changed and why, so that the framework's own evolution is auditable by design.

The final safeguard against a machine that agrees with you is a person who will not.